

AFL Assistant Minute Demo Script & Slide Outline

Slide 1: Title & Product Goal

On Slide

AFL Assistant

Domain-Scoped AI Assistant for AFL Statistics & Predictions


- AFL factual Q&A
- Structured data retrieval
- Match-winner prediction
- Domain guardrails
- Multi-turn conversation

Say

“Today I’m presenting my AFL Assistant, a domain-specific AI assistant built using LangGraph. The system is designed to answer AFL questions using structured data, provide match predictions using a machine-learning model, maintain conversation context, and refuse requests outside the AFL domain.”

Slide 2: Architecture

On Slide



Say

“The user interacts with the Streamlit interface. Requests are sent to a FastAPI backend, which passes them to the LangGraph controller. The graph applies guardrails and routes the request to direct answering, structured retrieval, prediction, or refusal. Retrieval uses AFL data, while prediction uses the trained match-winner model.”

Slide 3: Evaluation Results

On Slide

Category	Pass Rate
Factual Q&A	50%
Prediction Sanity	100%
Scope Guardrails	100%
Multi-turn	50%
Overall	73.3%

Prediction holdout: 67.13% accuracy

Say

“I evaluated the assistant using 30 test cases. Overall, 22 out of 30 passed, giving a 73.3 percent pass rate. Prediction sanity tests and scope guardrails both achieved 100 percent. The main weaknesses were natural-language retrieval mapping and multi-turn context handling. On the 2025 holdout set, the prediction model achieved 67.13 percent accuracy.”

Slide 4: Live Demo: Factual Question

Ask the assistant

What is Geelong’s record against Essendon?

Expected Result

Geelong Cats’ record against Essendon Bombers is 39 wins, 23 losses, and 1 draw across 63 matches.

Say

“This demonstrates grounded factual retrieval. Instead of asking the language model to remember the statistic, the system routes the request to the structured AFL retrieval tool.”

Slide 5: Live Demo: Prediction

Ask

Predict North Melbourne vs Fitzroy on 1983-04-04.

Expected Result

Predicted Winner:

North Melbourne Kangaroos

Probability:

72.0%

Date:

1983-04-04

Say

“This request takes a different path. LangGraph identifies it as a prediction request, the prediction tool resolves the fixture and sends the pre-match features to the trained model. The response includes the predicted winner and probability, along with a disclaimer that the prediction is not a certainty.”

Slide 6: Live Demo: Guardrail + Multi-turn

Test 1: Off-topic refusal

Ask:

Who won the latest cricket match?

Expected behavior:

“I can only assist with AFL-related questions.”

Say

“This demonstrates the domain guardrail. The assistant is intentionally restricted to AFL rather than becoming a general-purpose sports chatbot.”

Test 2: Multi-turn conversation

Ask:

What is Geelong’s record against Essendon?

Then:

How many wins did they have?

Expected answer:

39 wins.

Say

“The second test demonstrates conversation memory. The follow-up question does not repeat the teams, but the assistant uses the previous conversation context to understand that ‘they’ refers to Geelong.”

Slide 7: Limitations & Next Steps

On Slide

Current limitations

- Dataset requires regular updates
- Prediction accuracy is not perfect
- Some natural-language retrieval mappings fail
- Multi-turn edge cases remain
- Rule-based guardrails have edge cases

Next steps

- Improve query/tool mapping
- Strengthen structured conversation memory
- Weekly AFL data refresh
- Monitor prediction drift
- Retrain when justified by new data
- Deploy outside temporary Colab environment

Say

“The current system is functional, but there are clear areas for improvement. The most important next steps are improving natural-language tool selection, strengthening multi-turn context, keeping the AFL data current, monitoring prediction performance after each round, and eventually moving from the temporary Colab environment to persistent deployment.”

Closing: 15 seconds

“Overall, this project demonstrates an end-to-end domain-specific AI application combining LangGraph orchestration, structured retrieval, machine-learning prediction, guardrails, FastAPI, conversation memory, evaluation, and a user-facing interface.”